

11.1-Logistic Regression with Hyperparameter Tuning

August 3, 2025

1 Implementation Of Logistic Regression

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: ## creating the dataset using sklearn library
from sklearn.datasets import make_classification
```

```
[3]: ## Creating the dataset

X, y = make_classification(n_samples=1000, n_features=10, n_classes=2,
    random_state=42)
'''
So here:
n_samples is my total datapoint = 1000
n_features are the independent features or input = 10
n_classes = 2 (for binary)
'''
```

```
[3]: '\nSo here:\nn_samples is my total datapoint = 1000\nn_features are the
independent features or input = 10\nn_classes = 2 (for binary)\n\n'
```

```
[4]: X.shape, y.shape
```

```
[4]: ((1000, 10), (1000,))
```

```
[5]: df = pd.DataFrame(X)
```

```
[6]: df.head()
```

```
[6]:      0      1      2      3      4      5      6  \
0  0.964799 -0.066449  0.986768 -0.358079  0.997266  1.181890 -1.615679
```

```

1 -0.916511 -0.566395 -1.008614  0.831617 -1.176962  1.820544  1.752375
2 -0.109484 -0.432774 -0.457649  0.793818 -0.268646 -1.836360  1.239086
3  1.750412  2.023606  1.688159  0.006800 -1.607661  0.184741 -2.619427
4 -0.224726 -0.711303 -0.220778  0.117124  1.536061  0.597538  0.348645

```

```

          7          8          9
0 -1.210161 -0.628077  1.227274
1 -0.984534  0.363896  0.209470
2 -0.246383 -1.058145 -0.297376
3 -0.357445 -1.473127 -0.190039
4 -0.939156  0.175915  0.236224

```

```
[7]: df1 = pd.DataFrame(y)
```

```
[8]: df1[0].value_counts()
```

```

[8]: 0
      501
      499
      Name: count, dtype: int64

```

```

[9]: from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
      random_state=42)

```

```

[10]: ## Model training

      from sklearn.linear_model import LogisticRegression

      logistic = LogisticRegression()

```

```
[11]: logistic.fit(X, y)
```

```
[11]: LogisticRegression()
```

```

[12]: y_pred = logistic.predict(X_test)
      y_pred

```

```

[12]: array([0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0,
            1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0,
            0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1,
            0, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,
            0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1,
            1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0,
            1, 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1,
            0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1,
            0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1,
            1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 1,

```

```
1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1,
0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1,
0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1,
1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0])
```

```
[13]: # logistic.predict_proba(X_test)
```

```
[14]: ## Performance matrix
```

```
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report
```

```
[15]: score = accuracy_score(y_test, y_pred)
score
```

```
[15]: 0.8466666666666667
```

```
[16]: cm = confusion_matrix(y_test, y_pred)
cm
```

```
[16]: array([[116, 19],
[ 27, 138]], dtype=int64)
```

```
[17]: cr = classification_report(y_test, y_pred)
print(cr)
```

	precision	recall	f1-score	support
0	0.81	0.86	0.83	135
1	0.88	0.84	0.86	165
accuracy			0.85	300
macro avg	0.85	0.85	0.85	300
weighted avg	0.85	0.85	0.85	300

2 Hyperparameter Tuning And Cross Validation

```
[18]: model = LogisticRegression()

# penalty = ['l1', 'l2', 'elasticnet'] # here l1 is lasso regression, l2 is
# ridge regression and other is elasticnet regression
# c_values = [100, 10, 1, 0.1, 0.01]
# solver = ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']
```

```
[19]: ## Creating a dictionary for parameter tuning
# params = dict(penalty=penalty, C = c_values, solver=solver)
```

```

params = [
    # Solvers that only support l2
    {'solver': ['newton-cg', 'lbfgs', 'sag'], 'penalty': ['l2'], 'C': [100, 10, 1, 0.1, 0.01]},

    # liblinear supports l1 and l2
    {'solver': ['liblinear'], 'penalty': ['l1', 'l2'], 'C': [100, 10, 1, 0.1, 0.01]},

    # saga supports all penalties
    {'solver': ['saga'], 'penalty': ['l1', 'l2', 'elasticnet'], 'C': [100, 10, 1, 0.1, 0.01]}
]

```

3 GridSearchCV

```

[20]: from sklearn.model_selection import GridSearchCV, StratifiedKFold
      from sklearn.linear_model import LogisticRegression

      model = LogisticRegression()
      cv = StratifiedKFold(n_splits=5)
      grid = GridSearchCV(estimator=model, param_grid=params, scoring='accuracy',
                          cv=cv, n_jobs=-1)

      # Fit the model
      grid.fit(X_train, y_train)

      # Get the best score
      print(grid.best_score_)

```

0.8785714285714287

```
[21]: grid.best_params_
```

```
[21]: {'C': 0.01, 'penalty': 'l2', 'solver': 'newton-cg'}
```

```
[22]: y_pred = grid.predict(X_test)
```

```

[23]: score = accuracy_score(y_pred, y_test)
      print(score)
      print(classification_report(y_test, y_pred))
      print(confusion_matrix(y_pred, y_test))

```

0.8533333333333334

	precision	recall	f1-score	support
0	0.79	0.92	0.85	135

	1	0.92	0.80	0.86	165
accuracy				0.85	300
macro avg		0.86	0.86	0.85	300
weighted avg		0.86	0.85	0.85	300

```
[[124  33]
 [ 11 132]]
```

```
[24]: from sklearn.model_selection import cross_val_score
      from sklearn.metrics import accuracy_score

      best_model = grid.best_estimator_

      # Cross-validation accuracy on training data
      cv_scores = cross_val_score(best_model, X_train, y_train, cv=5)
      print("Mean CV Accuracy on Train:", np.mean(cv_scores))

      # Accuracy on test data
      test_accuracy = accuracy_score(y_test, best_model.predict(X_test))
      print("Test Accuracy:", test_accuracy)
```

Mean CV Accuracy on Train: 0.8785714285714287

Test Accuracy: 0.8533333333333334

Why the Scores Are Different:

grid.best_score_ = 87%: This is the average accuracy across 5-fold cross-validation on the training set only. It's used to find the best hyperparameters.

classification_report = 85%: After GridSearchCV found the best model, you tested it on new, unseen test data — naturally, performance can drop slightly here.

4 RandomizedSearchCV

```
[25]: from sklearn.model_selection import RandomizedSearchCV
```

```
[26]: rand = RandomizedSearchCV(estimator=model, param_distributions=params,
      ↪scoring='accuracy', cv=5, n_jobs=-1)
```

```
[27]: rand.fit(X_train, y_train)
```

```
[27]: RandomizedSearchCV(cv=5, estimator=LogisticRegression(), n_jobs=-1,
      param_distributions=[{'C': [100, 10, 1, 0.1, 0.01],
                           'penalty': ['l2'],
                           'solver': ['newton-cg', 'lbfgs',
                                       'sag']}],
      {'C': [100, 10, 1, 0.1, 0.01],
       'penalty': ['l1', 'l2']},
      scoring='accuracy', cv=5, n_jobs=-1)
```

```

        'solver': ['liblinear']},
        {'C': [100, 10, 1, 0.1, 0.01],
         'penalty': ['l1', 'l2', 'elasticnet'],
         'solver': ['saga']}],

        scoring='accuracy')

```

```
[28]: rand.best_estimator_
```

```
[28]: LogisticRegression(C=0.1, penalty='l1', solver='liblinear')
```

```
[29]: rand.best_score_
```

```
[29]: 0.8742857142857143
```

```
[30]: y_pred = rand.predict(X_test)
score = accuracy_score(y_pred, y_test)
print(score)
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_pred, y_test))
```

```
0.8533333333333334
```

	precision	recall	f1-score	support
0	0.79	0.91	0.85	135
1	0.92	0.81	0.86	165
accuracy			0.85	300
macro avg	0.86	0.86	0.85	300
weighted avg	0.86	0.85	0.85	300

```
[[123  32]
 [ 12 133]]
```

Same accuracy difference as GridSearchCV